


Comparing AI Agents Performance in 2048

By Lucas Summers, Braeden Alonge, Nathan Lim, and Megan Fung




The Team: Contributions & Collaboration



Megan Fung

Computer Science '26

- Heuristics, Expectimax
- Presentation
- Report



Braeden Alonge

Computer Science '26

- Expectimax algorithm
- Report
- Presentation



Lucas Summers

Computer Science '26

- MCTS + Hybrid Models
- Heuristic design
- Flask server + applet
- Report + Presentation



Nathan Lim

Computer Science '26

- Reinforcement Learning agent / training
- MCTS + RL Hybrid Agent
- Report
- Presentation

Abstract

GOAL Build high performing AI agents to play 2048.

AI Methods Expectimax, MCTS, Reinforcement Learning (DQN), Hybrid

Outcome Evaluate agent performance through simulations and heuristics.

Purpose Benchmark planning algorithms in a stochastic environment.

Motivation

1

2048 is simple to play, hard to master, and ideal for AI experimentation.

2

Combines planning, randomness, and strategy.

3

Demonstrates real world decision-making under uncertainty.

4

Helps evaluate how AI agents can adapt and optimize gameplay over time.

Background + Prior Work

Expectimax:

- Used in many 2048 bots for modeling random tile placement
- Effective but struggles with large branching factors

MCTS:

- Balances exploration and exploitation.
- Uses a statistical framework to guide decision-making.

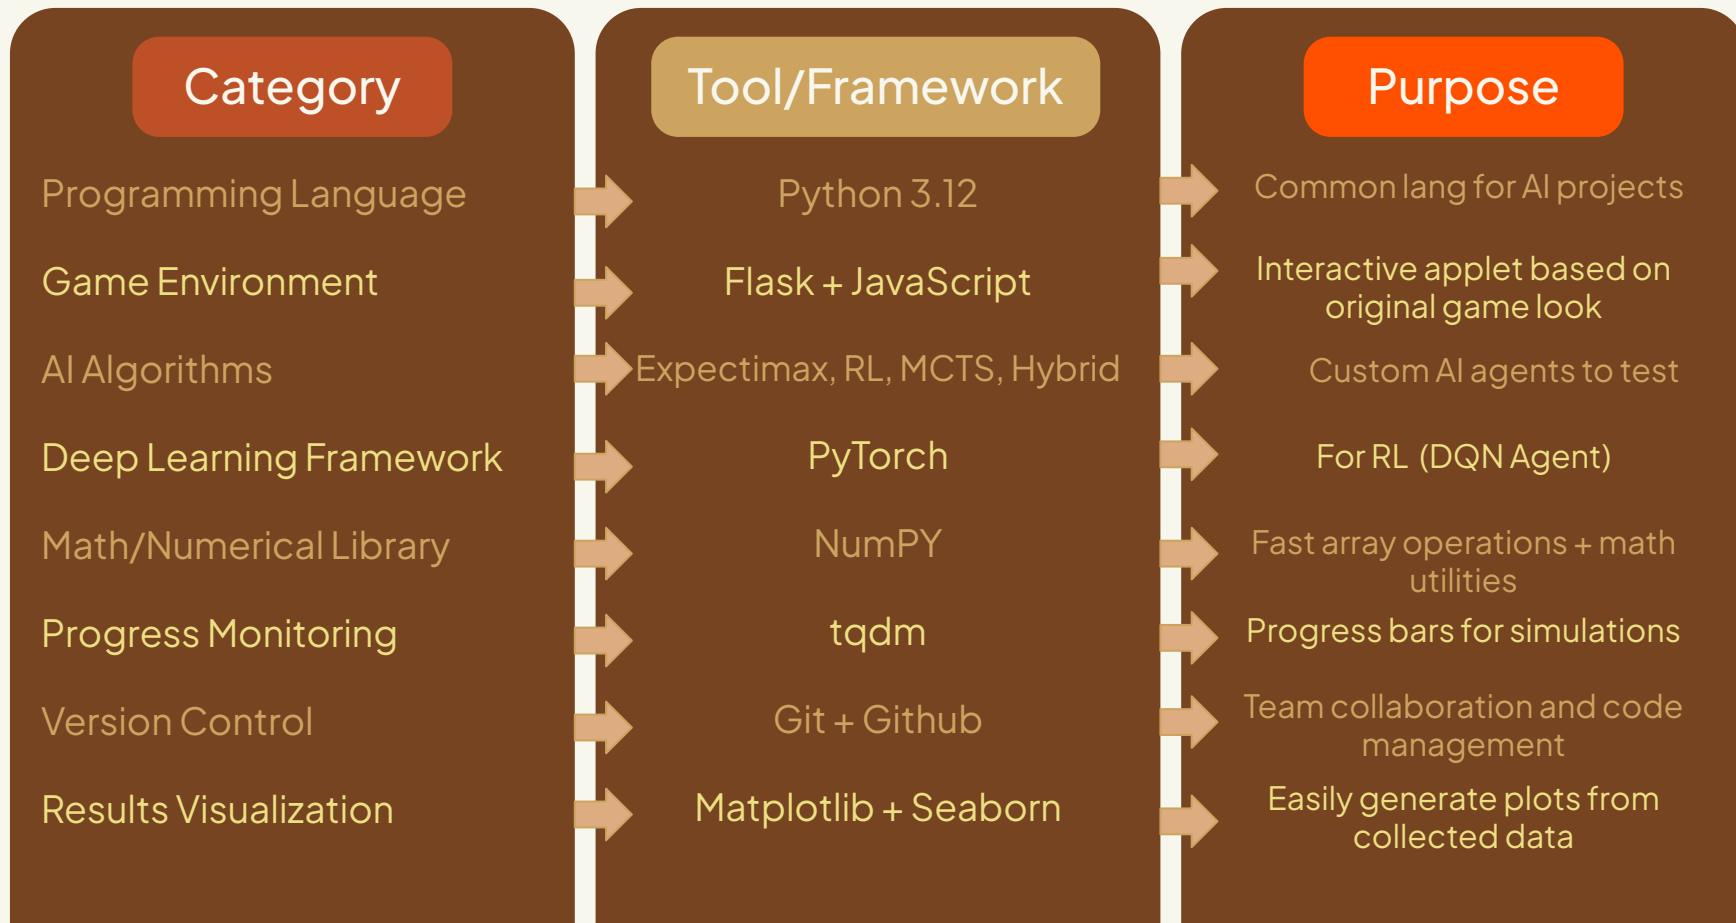
Reinforcement Learning:

- Requires significant training.
- Learn optimal strategies from scratch through self play.

Hybrid:

- Combines strengths of MCTS with predictive evaluation
- Better quality rollouts leads to better decisions overall

Toolchain & Technical Stack




Overview of AI Agents

- **Random:** Baseline, no planning
- **Greedy:** Chooses moves based on immediate rewards
- **Expectimax:** Explores deterministic moves and stochastic tile spawns
- **MCTS:** Simulates outcomes using random rollouts and UCB1
- **Reinforcement Learning:** Learns state action values (Q values) from gameplay
- **Hybrid (w/ Expectimax):** MCTS + Expectimax rollouts for smarter simulations
- **Hybrid (w/ RL):** MCTS + RL rollouts

2048 AI!

SCORE
204

BEST
860



Stop AI

New Game

AI Agent:

- ☐ Random ☐ Greedy ☒ Mcts
☐ Hybrid ☐ Expect ☐ RI ☐ Mcts_rl

Speed:

- ☐ Very Fast ☒ Fast ☐ Medium ☐ Slow

Agent Statistics:

Thinking Time:	0.50
Avg Reward:	-260.24
Execution Time:	0.50
Max Depth Reached:	12
Search Iterations:	95

Highest numbers:



Heuristics

We use a weighted composite of heuristics to evaluate how promising a board is for achieving high scores. These heuristics guide agents like Expectimax and MCTS, whose game tree rarely reaches a terminal state.

Empty Tile

- Rewards boards with more empty spaces
- Empty tiles → greater flexibility for future moves

Monotonicity

- Encourages tile ordering in increasing/decreasing sequences along rows + columns

Smoothness

- Measures how similar adjacent tiles are:
- Differ in large amounts → penalized
- Gradual value transitions → rewarded

Quality/Stability

- How resistant the board is to quality loss
- Rewards the max tiles being protected by other high tiles

Corner Max

- Encourages strategic placement of the highest tile in a corner, a common expert strategy

Merge Potential

- Evaluates how many adjacent tiles are able to be merged (either vertically or horizontally)

Experiment Design

Agents Tested:

- Random
- Greedy
- Expectimax
- MCTS
- Reinforcement Learning
- MCTS/Expectimax Hybrid
- MCTS/RL Hybrid

Each agent plays 100+ games with a fixed random seed for reproducibility.

Avg, Min, and Max Final Score

Win Rate

Average Moves Per Game

Efficiency (score/moves)

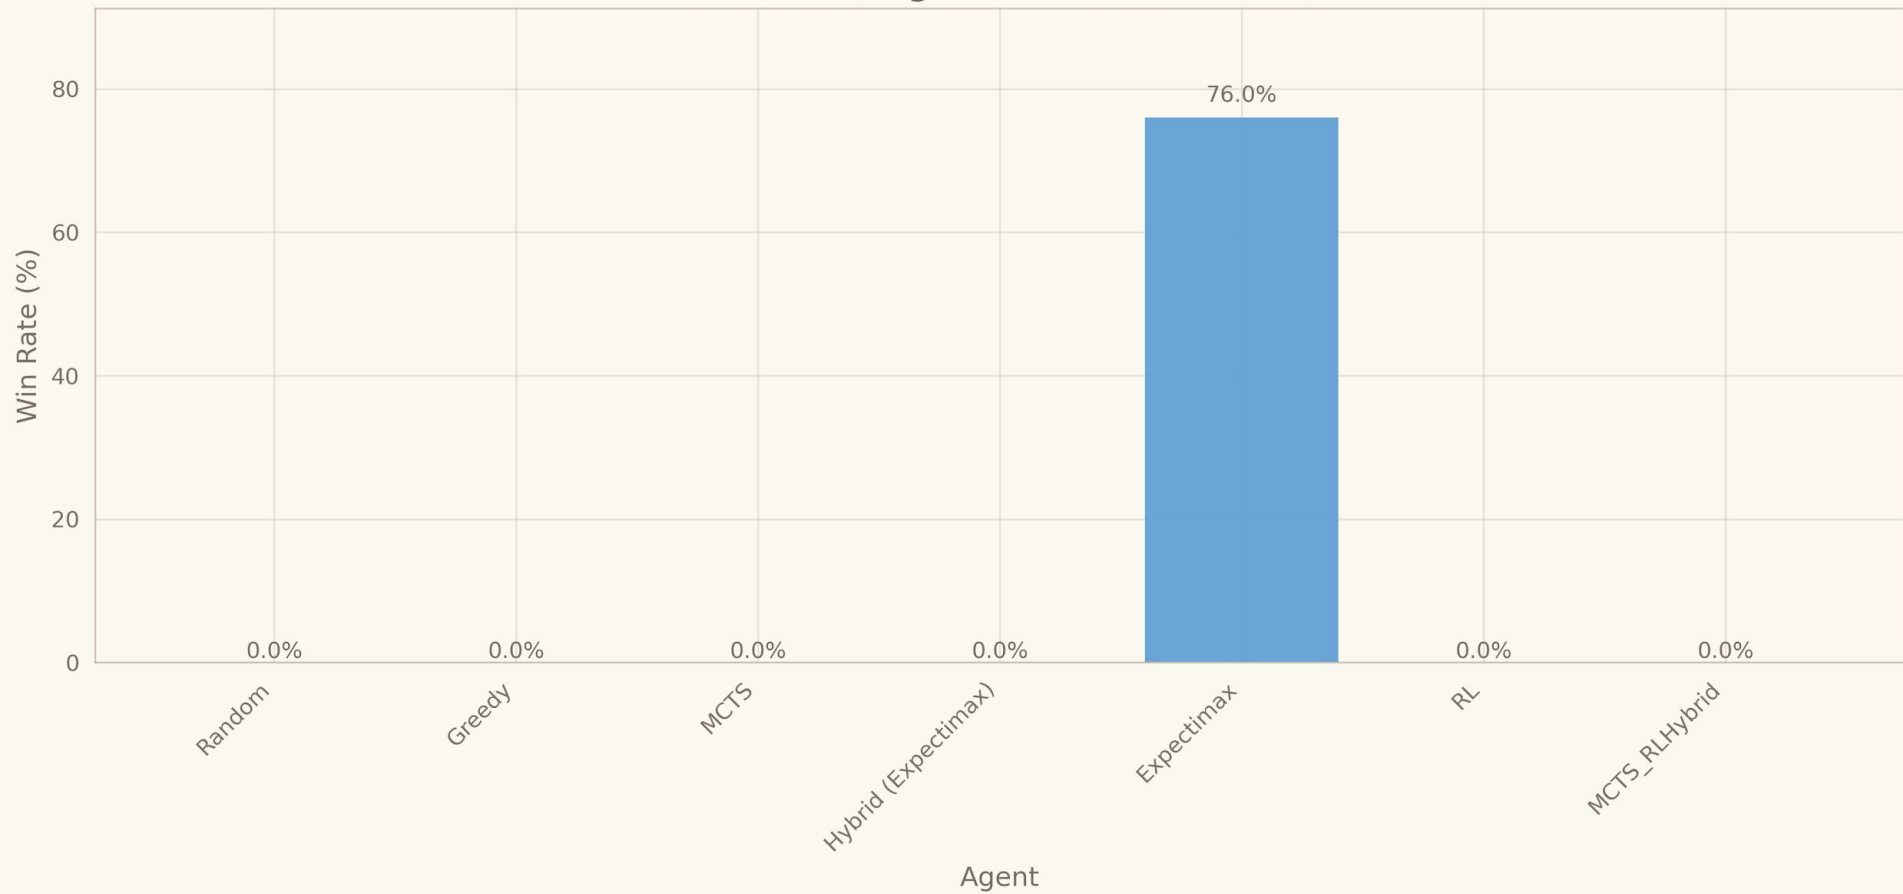
Median Max Tile

Absolute Max Tile

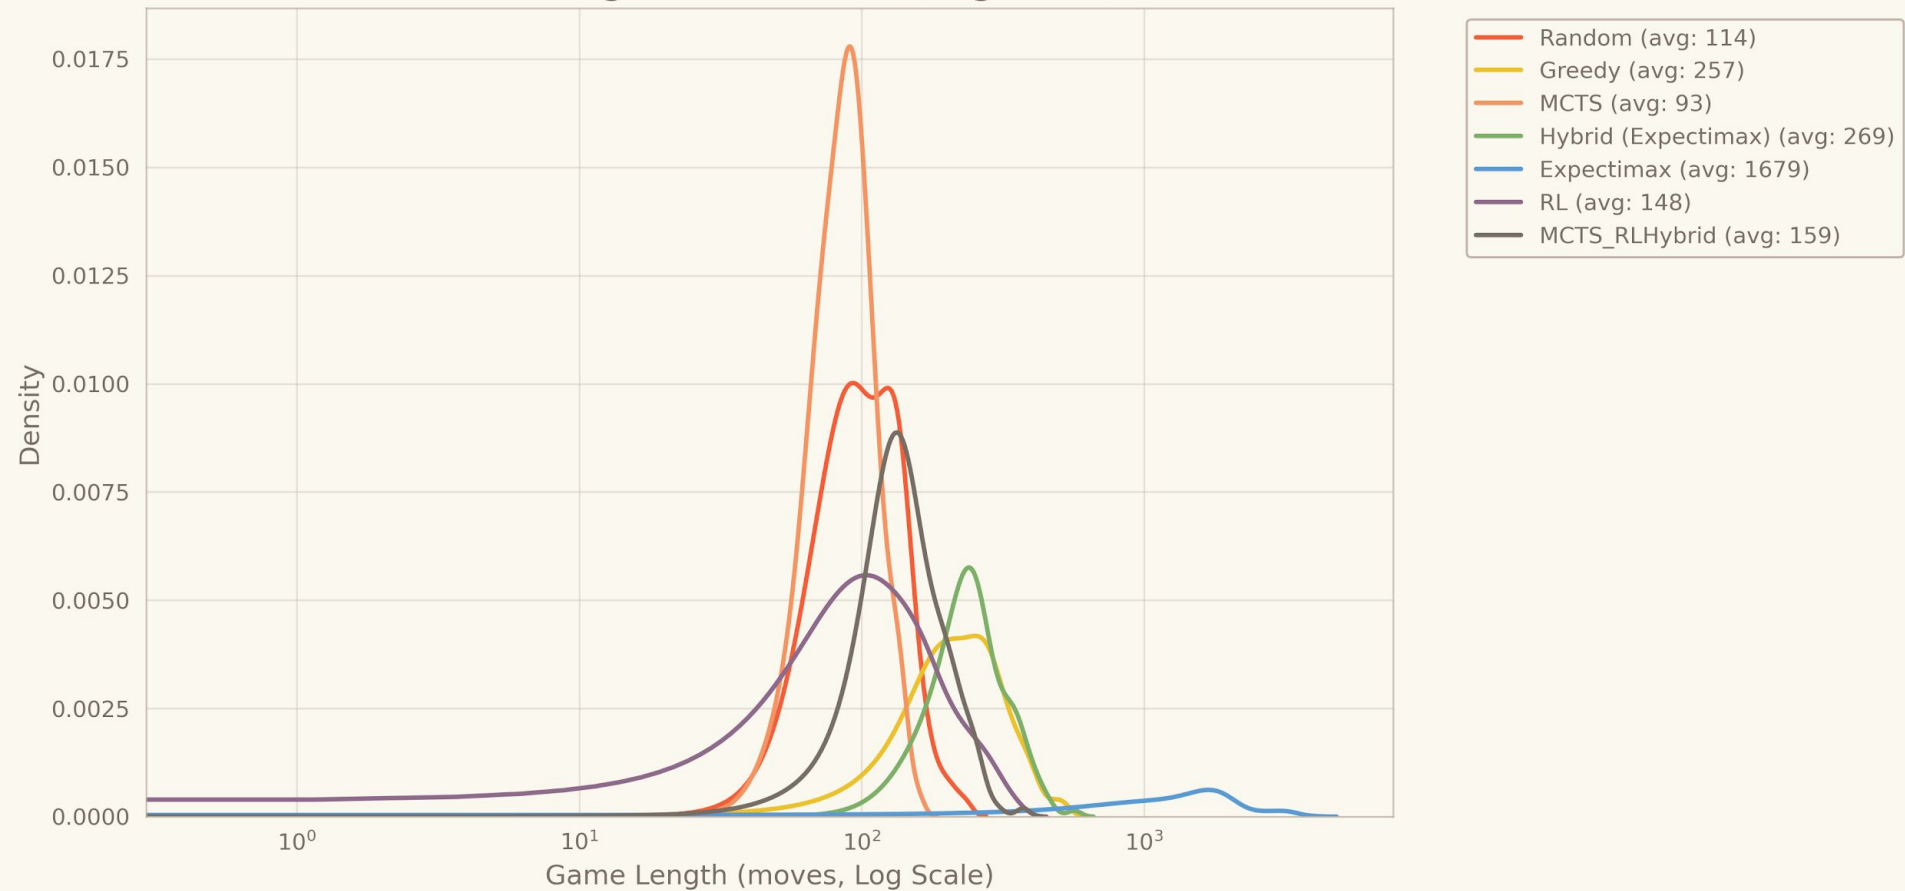
Demo




Win Rate (% games with 2048+ tile)




Game Length Distributions (Log Scale)




Score Distributions (Log Scale)




Efficiency (Score per Move)



Max Tile Distributions (Highest 9 Tiles)



Move Distributions



Results Conclusion

1. **Expectimax** achieved the highest absolute max tile (4096) and average score (33490.32), outperforming all other agents
2. **MCTS + Expectimax Hybrid** agent follows with 5081.01 average score and decent tile depth
3. **Greedy** agent has surprisingly good performance, coming in third
4. **MCTS** (Random), **MCTS** (RL), and **RL** showed balanced performance, with RL slightly outperforming MCTS in score
5. **Random** agent, as expected, had the lowest score and tile depth

These trends validate the strength of structured planning and evaluation in Expectimax and show Greedy agents can perform well with minimal computation.

Discussion

Observations

- Structured planning outperforms reactive or random strategies
- Expectimax's lookahead pays off despite longer computation time
- RL showed promise with generalization, though training time limited its performance

Challenges

- Runtime efficiency for deeper search agents
- RL tuning and training resource constraints
- Difficulty comparing agents with different decision speeds

Future Work

- Train deeper and more complex RL models
 - Double DQN or Dueling Networks)
- Explore more rollout strategies for MCTS
- Integrate AlphaZero-style RL
- Expand the analysis to larger board sizes (ex. 5×5 or 6×6) and compare scalability
- Evaluate computational efficiency more rigorously
 - Measure time per move under different hardware configurations
- Test adding, removing, or modifying existing heuristics
- Implement a grid search for ideal hyperparameters

Limitations

Thinking time:

- Less predictions and evaluation allowed for each algorithm
- Less thinking time means less effectiveness

Computational Power:

- No access to NVIDIA GPUs for CUDA
- Reduced efficiency and tasks completed with the given limited thinking time

Time Constraint:

- Scope of the project was too large to produce desired results
- More time would allow for more research and better development of possible algorithms and heuristics

Stochastic Game:

- High variance led to poor algorithm effectiveness
- RL and MCTS struggled the most

Retrospective

Successes

- Built a modular system with swappable agents and visual feedback
- Designed and implemented competitive AI agents using diverse strategies
- Created reusable evaluation metrics and tools

Lessons

- Heuristics matter immensely in search based agents
- Agent design is a balance between performance and practicality
- RL requires large amounts of training time and computational power

Reflection

- Developed a deeper understanding of reinforcement learning's real-world limitations
- Recognized the trade-offs between model accuracy and computational cost/thinking time
- Learned the importance of modular design for scalability and future improvements

Conclusion

→ AI planning must balance depth, adaptability, and efficiency

Expectimax: High accuracy, low speed

MCTS: Balanced performance

RL: Potential for learning beyond heuristics

Hybrid: Promising results with smart rollouts

2	2	8	2
	8	16	2
128	64	32	4
256	512	128	16

Thank
You!

Q & A...